

3-ID-C BITS

A tutorial walkthrough

Self-paced; ~30 slides; ~45-60 minutes

Presented 2026-06-01

Audience

This deck is for someone who:

- Knows the 3-ID-C hardware and beamline workflow
- Has used SPEC or bare EPICS (`caget` / `caput` /MEDM)
- Has *some* Python familiarity (knows what a dict and a function are; not necessarily comfortable writing classes)
- Has heard "Bluesky is replacing SPEC" and wants to know what that actually means in practice

If you've already seen the standard intro deck, this one fills in the *why*.

What we'll cover

1. The shape of a BITS session
2. SPEC -> Bluesky: command cross-walk
3. EPICS -> ophyd: the device model
4. Plans, plan stubs, and `yield from`
5. The RunEngine: pause/resume, suspenders, document streams
6. What's in the 3-ID-C instrument today
7. The omega <-> laser_optics interlock (design + scope)
8. Inspecting past data via Tiled
9. The docs site and how to extend it
10. Where Bluesky is worse than SPEC (honest section)

Part 1: shape of a session

```
conda activate 3idc-bits  
ipython
```

```
from id3c.startup import *
```

That `from ... import *` triggers:

- Load `iconfig.yml` (top-level session config)
- Create the **RunEngine** (`RE`) with `bec` and `tilde` subscribers
- Instantiate every device in `configs/devices.yml`
- Install the omega <-> laser_optics interlock
- Register IPython magics (`%wa` , `%mov` , `%ct`)
- Import demo plans (`sim_print_plan` , ...)

Part 1: what's bound at the prompt (1/2) -- machinery

```
RE          # the RunEngine
oregistry   # the device registry
cat         # the Tiled catalog client
bec         # BestEffortCallback (live plots/tables)

bps         # bluesky.plan_stubs (mv, sleep, ...)
bp          # bluesky.plans (count, scan, ...)
```

Part 1: what's bound at the prompt (2/2) -- devices

```
sample_stage      # a MotorBundle
laser_optics      # our custom LaserOptics
shutter           # ApsPssShutter
eiger2            # area detector
sim_motor, sim_det # simulators
```

`%wa` shows everything. Tab-completion works on devices.

`listobjects()` Concise table of all top-level control devices.

Part 2: SPEC -> Bluesky (1/2) -- motion and inspection

SPEC	Bluesky
<code>mv samx 5</code>	<code>RE(bps.mv(sample_stage.xprime, 5))</code>
<code>mvr samx 0.1</code>	<code>RE(bps.mvr(sample_stage.xprime, 0.1))</code>
<code>wm samx</code>	<code>sample_stage.xprime.position</code> (no <code>RE</code> !)
<code>wa</code>	<code>%wa</code>
<code>shopen</code>	<code>RE(bps.mv(shutter, "open"))</code>

Part 2: SPEC -> Bluesky (2/2) -- counts and scans

SPEC	Bluesky
ct 1	RE(bp.count([scaler]))
ascan samx 0 10 10 1	RE(bp.scan([scaler], sample_stage.xprime, 0, 10, 11))
mesh ...	RE(bp.grid_scan([scaler], mot1, ..., mot2, ...))

Note SPEC's "10 intervals" vs. Bluesky's "11 points".

Part 2: things SPEC does that Bluesky doesn't

Being honest:

- **Compactness.** `ascan samx 0 10 10 1` is shorter than the Bluesky equivalent. We can alias common commands, but bare commands are longer.
- **Macros.** `overnight.mac` is faster to write than authoring a Python plan.
- **One command -> one file.** SPEC files are human-readable text. Bluesky runs live in a Tiled catalog; you need a client to read.
- **Decades of stability.** SPEC's command set is unchanged; Bluesky is younger.
- **Programming Language --** SPEC's uses a familiar C-like syntax but is unique to SPEC, the language support is limited to the SPEC user community.

Part 2: things Bluesky does that SPEC doesn't (1/2)

- **Structured metadata** -- every run has a UID, a `scan_id`, and a `md=` dict you can search later.
- **Pause / resume.** `ctrl-c ctrl-c` pauses mid-scan; `RE.resume()` continues.
- **Document streams** -- `start`, `descriptor`, `event`, `stop` documents flow to *subscribers* (BEC for live plots, TiledWriter for storage). Standardized format.
- **Suspenders.** "Pause on beam dump, resume when it returns" is a generic mechanism, not bespoke per-beamline code.

Part 2: things Bluesky does that SPEC doesn't (2/2)

- **Programming Language** -- Python is a popular & well-documented language. Help is available from many sources.
- **Syntax Checking** -- Syntax checking is inherent to Python. SPEC macros are string text, interpreted at run time.
- **Catalog-backed history.** `cat[-1]` is the most recent run, `cat[uid]` is a specific one; you can search by metadata.
- **Area detectors.** ophyd wraps the full EPICS areaDetector framework: the cam, the plugin chain (ROI, stats, codecs, file writers, PVA push, ...), and per-run orchestration of all of it. HDF5-via-external-links is one example; see [how_to/visualize_hdf5](#).

Part 3: EPICS -> ophyd

You used to type the PV string:

```
caget 3idxps1:m5.RBV  
caput 3idxps1:m5.VAL 30
```

Now you address a Python object that wraps the PVs:

```
sample_stage.omega.user_readback.get()    # the .RBV  
sample_stage.omega.user_setpoint.put(30)  # the .VAL  
sample_stage.omega.move(30)               # set + wait
```

An `EpicsMotor` wraps ~12 PVs (.VAL, .RBV, .DMOV, .MOVN, .STOP, .HLM, .LLM, .EGU, .OFF, ...). You access them as attributes.

Part 3: why wrap PVs?

The trade:

- You **lose** "any PV at any time" -- you have to define what signals a device has.
- You **gain**:
 - Self-documenting object (tab-completion!)
 - Search PV and connect once.
 - Long-lived connection with cached reads
 - Subscription-first API
 - `read()` -> structured dict ready for archiving
 - Integration with the Bluesky document stream

For *one* CA operation, `caget` is still fine. For an **instrument**, the wrapper pays for itself many times over.

Part 3: `get()` vs `read()`

Two operations users confuse:

```
sample_stage.omega.user_readback.get()  
# 30.0                -- ONE signal's value  
  
sample_stage.omega.read()  
# {'sample_stage_omega': {'value': 30.0, 'timestamp': ...},  
#  'sample_stage_omega_user_setpoint': {'value': 30.0, 'timestamp': ...}}  
# -- ALL signals of kind hinted/normal
```

`get()` for "give me a number." `read()` for "give me a snapshot."

`device.read()` is what the RunEngine calls internally during a scan.

Part 4: plans and plan stubs

A **plan** publishes Bluesky documents: examples `bp.count` , `bp.scan`

A **plan stub** does not publish: examples `bps.mv` , `bps.sleep`

Both are *generators*: functions that yield messages for the RunEngine. Both work with `RE(...)` .

The difference matters when you *write* plans:

- Plan stubs are easy: yield from other stubs, do not bracket a run.
- Plans need `open_run` / `close_run` (or compose a `bp.*`).

Most user code is plan stubs. Composing into a plan is usually wrapping an existing `bp.*` plan and including calls to plan stubs as needed.

Part 4: `yield from`

`yield from` is the Python syntax for **composing one generator inside another**.

```
@plan
def my_plan():
    yield from bps.mv(motor, 5)      # composed stub
    yield from bp.count([detector])  # composed plan
```

You use it **inside** a plan you are writing. You do not use it at the IPython prompt. At the prompt, use `RE(...)`.

If you forget either, the generator is created and discarded -- the call does nothing. The `@plan` decorator catches that.

Part 4: the `@plan` decorator

All plans and plan stubs we author are decorated with `bluesky.utils.plan`. If you call one without `RE(...)`, a warning prints shortly after you press Enter:

```
RuntimeWarning: plan `sim_print_plan` was never iterated,  
               did you mean to use `yield from`?
```

The warning's traceback points at *your* command line, not at internal Bluesky code, so you can see exactly which line to retype.

Convention: every new plan/plan-stub in `src/id3c/` gets `@plan`. See `AGENTS.md` > "`@plan` decorator on our own plans".

Part 5: the RunEngine

The RunEngine is the thing that **executes** a plan. It:

- Iterates the generator one message at a time
- Dispatches each message to the appropriate device
- Publishes documents to subscribers (BEC, TiledWriter)
- Receives CA monitors and caches values
- Handles pauses, suspenders, errors, cleanup
- Threads metadata through the document stream

You can think of it as the SPEC interpreter, but for plans.

The RunEngine is *the* thing that turns a description of a scan into an actual scan.

Part 5: pause / resume

During a long scan:

- `ctrl-c` once -- deferred pause (after current message finishes)
- `ctrl-c ctrl-c` -- immediate pause (interrupts the current await)

At the pause prompt:

```
RE.resume()      # continue
RE.abort()       # finish, exit_status='abort'

# used less often
RE.stop()        # finish cleanly, success
RE.halt()        # emergency stop, no documents
```

This is **free** -- works for every plan including custom ones.

Part 5: subscribers and the document stream

Every `RE(plan)` invocation emits a stream of documents:

document type	when
<code>start</code>	once, at plan begin
<code>descriptor</code>	once per data stream
<code>event</code>	once per data point
<code>stop</code>	once, at plan end

Subscribers consume the stream live:

- `bec` -- `BestEffortCallback`: prints tables, opens plots
- **TiledWriter** -- sends documents to the Tiled server
- **nxwriter** (optional) -- writes NeXus-format HDF5 files

You can add your own: `RE.subscribe(my_callback)` .

Part 6: what's installed today

device	notes
sample_stage	xprime / base_y / zprime / omega (interlocked)
detector_stage	det_x / eiger_y / eiger_z
laser_optics	us / ds (interlocked with omega)
shutter	A-station PSS shutter
eiger2	Eiger2 500k; HDF5 file plugin still FIXME
sim_motor , sim_det	simulators for verification

`%wa` lists everything by label; `%wa baseline` shows the devices recorded at the start and end of every run.

Part 7: the omega <-> laser_optics interlock

Why: when the laser optics are not retracted, they could collide with the rotating sample stage; symmetrically, pulling the laser in/out while omega is moving is also risky.

What is protected:

- `sample_stage.omega` blocked unless `laser_optics.is_out` (both axes within +/- 1 mm of -75 mm)
- `laser_optics.us / .ds` blocked while `omega.motor_is_moving`

How: `InterlockedEpicsMotor.move()` runs an interlock check before any CA put (pre-flight) **and** subscribes a watcher to the relevant signals during motion (mid-flight). Failure raises `MotionInterlock`.

Part 7: scope of the interlock

This is a Python-session interlock. It does not:

- Write to EPICS PV disable fields
- Install IOC sequencer code
- Protect against MEDM jogs
- Protect against `caput` from a shell
- Protect against a different Bluesky session
- Survive a Python process crash

For *session-independent* hardware-grade protection, the right place is the IOC (CALC/SCALC, state notation, or a soft record driving PV disable). Adding that is a separate (welcome) project.

This is documented honestly in the module docstring of `interlocked_motor.py` and in `docs/source/explanation/interlocks.md`.

Part 8: inspecting past data

Bluesky runs are written to a [Tiled](#) server (3-ID-C uses <http://sn.xray.aps.anl.gov:8000>). The session-level client is `cat` :

```
cat[-1]                # most recent run
cat["<uid>"]           # by UID

run = cat[-1]
run.metadata["start"]  # plan args, scan_id, plan_name, ...
run.primary.read()     # xarray Dataset of the main stream
run.baseline.read()    # baseline-labeled devices
```

Part 8: image data and HDF5 external links

For an area-detector run, the data flow is:

```
Eiger IOC -> writes image.h5  
custom Bluesky plan -> writes master.h5 with HDF5 external link to image.h5  
TiledWriter -> sends run docs referencing master.h5  
client reads: client -> Tiled -> master.h5 -> image.h5 (via link)
```

The chain works *if* every hop succeeds, especially the last (image file visible to the Tiled server). At 3-ID-C, this end-to-end path is **not yet validated**. See

`docs/source/how_to/visualize_hdf5.md` for the current state.

Part 9: the docs site

`docs/source/` is a Sphinx site, [Diátaxis](#)-organized:

- **tutorials/** -- *learning* (first session, SPEC->Bluesky, EPICS->ophyd)
- **how_to/** -- *task* (add a device, add a plan, inspect data, ...)
- **reference/** -- *lookup* (cheat sheet, quick reference, configuration)
- **explanation/** -- *understanding* (RunEngine, plans+stubs, interlocks)

Plus this presentations directory.

Build locally: `cd docs && make html`. CI deploys `main` to `https://bcda-aps.github.io/3idc-bits/`.

Part 9: how to extend the docs

Add a...	...where
New page	<code>docs/source/<section>/<name>.md</code> + toctree entry
New device	<code>src/id3c/configs/devices.yml</code> (or a class in <code>src/id3c/devices/</code>)
Custom motor	<code>mb_creator</code> per-axis <code>class:</code> key in YAML
Plan	<code>src/id3c/plans/<topic>.py</code> + <code>startup.py</code> import
Interlock	<code>src/id3c/devices/<a>__interlock.py</code> + <code>startup.py</code> line
Subscriber callback	<code>src/id3c/callbacks/<name>.py</code> + <code>RE.subscribe</code>

Every category has a matching how-to page; the `reference/quick_reference.md` table has the full mapping.

Part 10: honest summary

Bluesky gains (vs SPEC):

- Reproducibility, recoverability, post-experiment data access
- Pause / resume, suspenders, structured metadata
- Live plots and tables for free
- A real software stack you can hire Python developers for

Bluesky pains (vs SPEC):

- More verbose syntax
- A learning curve (this deck exists for a reason)
- More layers between you and the PVs when things go wrong
- Tracebacks are long; learn to read the bottom line

A worthy trade to gain recoverable data and reproducible workflows.

Where to go from here

- Open IPython. Run `from id3c.startup import *`. Type `%wa`.
- Run the three sim plans: `RE(sim_print_plan())`, `RE(sim_count_plan())`, `RE(sim_rel_scan_plan())`. Watch BEC plot the third one.
- Try forgetting `RE(...)` once on purpose: `sim_print_plan()`. See the `RuntimeWarning`. Internalize it.
- Read the cheat sheet. Print it. Tape it next to your monitor.
- File issues for anything that confuses you.

Welcome to Bluesky at 3-ID-C.

References

- This repo: <https://github.com/BCDA-APS/3idc-bits>
- Docs: <https://bcda-aps.github.io/3idc-bits/>
- Bluesky upstream: <https://blueskyproject.io/>
- apsbits: <https://github.com/BCDA-APS/apsbits>
- Tiled: <https://blueskyproject.io/tiled/>
- Diátaxis (doc structure): <https://diataxis.fr/>
- **Bluesky Office Hours:** [Every Wednesday, 2-3 pm on Teams](#)

Questions: <https://github.com/BCDA-APS/3idc-bits/issues>